

# numpy 1D arrays

› go over solutions to homework 2, and compare by reference vs. value from last class

› numpy

- numpy is the main numerical library in Python, basis for many other scientific Python libraries
  - typical usage: `import numpy as np`
  - numpy provides: 1. the `ndarray` object, 2. lots of numerical and array functions
  - arrays are sequences, like lists and tuples, but for large datasets are faster and much more memory efficient
  - unlike lists, can explicitly be multidimensional - useful for e.g. images and movies
  - only deal with 1D for now
  - tradeoff: not as flexible as lists - for efficiency, each entry in an array has to be of the same data type
    - you can have an array of ints, or floats, or strings or booleans, but not a mixture
    - so far, we've seen that there are two main numeric data types: int and float
    - there are also different kinds of integer and float data subtypes (see class 05), each array can contain only one kind
  - like a tuple, array length generally **can't** change, but like a list, its values **can** be changed, so it's "semi-mutable"
- lots of ways to initialize an array
  - explicitly: first create a list or a tuple, then convert to an array:
  - `a = np.array([1, 2, 3])` or `a = np.array((1, 2, 3))`
  - `a = np.arange(10)` returns a range of integers
    - very similar to `list(range(10))`, but returns an array instead of a list
  - `a = np.zeros(10)` - an array with 10 entries, all `0.0`
  - `a = np.ones(10)` - an array with 10 entries, all `1.0`
  - `a = np.random.random(10)` - 10 random numbers uniformly distributed between 0 and 1
  - `a = np.tile(5, 10)` - 10 copies of the integer 5
  - `a = np.tile([1, 2], 5)` - 5 copies of the sequence `[1, 2]`
  - `a.fill(7)` fills the existing array `a` with the number 7
  - array methods (e.g. `a.fill()`) usually operate on the array in-place, while numpy functions (e.g. `np.zeros()`) usually return a new array
  - here's an exception: `b = a.copy()`
  - numpy functions often have array method counterparts (and vice versa)
    - `copy()` and `sort()` are two examples:

```
a = np.random.random(10)
b = a.copy()
c = np.copy(b)
```

    - are `a`, `b` and `c` equal? test with `==`, get a boolean answer for each entry
    - are `a`, `b` and `c` the same objects? test with `is`, get a single bool answer

```
d = a
d.sort() # in-place
e = np.sort(b)
```

- are `a` , `d` and `e` equal? are they the same objects?
  - are `b` and `e` equal? are they the same object?
  - can use `id()` to check the unique memory address of an object
- like other sequences (tuples & lists), get length of array using `len(a)` , but can also get array shape using the `.shape` attribute
    - `shape` returns the length along all dimensions of `a`
    - length of the first dimension is `a.shape[0]` , identical to `len(a)`
    - `a.ndim` tells you the number of dimensions - multidimensional arrays covered later
  - indexing in 1D is the same as for tuples & lists: 0-based, -ve indices count from the end
    - `a[0]` = 7 assigns 7 to 1st entry
    - `a[1]` = 7 assigns 7 to 2nd entry
    - `a[-1]` = 7 assigns 7 to last entry
    - `a[-2]` = 7 assigns 7 to 2nd last entry
  - slicing in 1D is also the same as for tuples and lists
    - retrieve a slice: the first 5 entries
      - `b = a[0:5]` or `b = a[:5]`
  - unlike lists, with arrays, you can also *assign* values to a slice:
    - assign to the last 5 entries
      - `a[5:10] = 7` or `a[5:] = 7`
    - assign to all entries with `:` , i.e. slice from start to end
      - `a[:] = 8` , same as `a.fill(8)`
    - what happens if you do `a = 8` ?
  - arrays also have "fancy" indexing:
    - allow you to ask for multiple values from an array at once
    - two types: **integer** & **boolean** fancy indexing
    - both are kind of a hybrid between normal indexing and slicing
      - benefit of fancy indexing over slicing is that you can specify any sequence of indices, not just evenly spaced ones
      - you can even specify the same index multiple times
    - integer fancy indexing
 

```
a = np.random.random(10) # init an array of random data
i = [3, 7, 5, 2, 7] # create a list of indices
vals = a[i] # index into array using integer fancy indexing
a[i] = -1 # assign -1 at multiple locations using integer fancy indexing
```

      - can ask for array values in arbitrary order

- can ask for the same value repeatedly
- can't do integer fancy indexing with lists:

```
l = list(range(10))
l[i] # TypeError: list indices must be integers or slices, not list
```

- boolean fancy indexing
  - ask a question of values of the array, get an answer back made up of boolean values of same length as original array
  - `i = a >= 0.5` returns an array of booleans, which can be used for indexing
  - `a[i]` or `a[a >= 0.5]` returns only those entries in `a` that are `>= 0.5`
  - i.e., where `i` is `True`, return the value in `a` at that index
  - what if you have another array `b` that is of different length? can you also index into it with the above `i`?

```
b = np.random.random(3)
b[i] # IndexError
```

- again, can't do boolean fancy indexing with lists: `l[i] # TypeError`

- **vectorized** math operators (`+`, `-`, `*`, `/`, `**`) and comparitors (`==`, `>`, `>=`, `<`, `<=`, `!=`)

- vectorized: work on all values of an array at the same time
- `a = np.array([1, 2, 3])`
- arrays & scalars
  - `a + 1` returns a new array with 1 added to all the entries in `a`
  - `a += 1` increments all entries in `a` in-place by 1, doesn't return anything
  - `a -= 1` decrements all entries in `a` in-place by 1, doesn't return anything
- `b = np.array([4, 5, 6])`
- `a + b` returns another array whose values are the sum of the corresponding two values in `a` and `b`
  - in comparison, what does `+` do for strings? for lists?
  - use `np.concatenate([a, b])` or `np.concatenate((a, b))` to combine arrays sequentially
- what happens if you try to do one of the above vectorized operations on two arrays of different length?

## › array exercises:

1. Use a for loop to build a list of 3 arrays, each array of length 5, initialized to zeros
2. Find the vector difference between the following two arrays and assign it to a new array called `d`:

```
a = np.array([10, 20, 30, 40, 50])
b = np.array([5, 12, 18, 31, 45])
```

3. "Reduce" `d` to a single number by using the function `np.mean()` or the method `d.mean()`
4. Write a function called `rms()` that calculates the RMS (root mean square) of an input array. RMS is the the square root of the mean of the square of a signal. To calculate square root, use the function `np.sqrt()`. RMS can be calculated in a single line.
5. Use your `rms()` function to calculate the RMS of the difference between the two arrays in 2.
6. Concatenate `a` and `b` into a new array called `c`. Now sort `c`, either using the function `np.sort()` or the method `c.sort()`
7. Create a boolean array `i` that describes where the values in `c` fall between 10 and 20 (inclusive). Hint: use the `*` or `&` operators to perform a vectorized AND operation between two boolean arrays.
8. Use `i` to extract the corresponding values from `c`
9. Use *integer* fancy indexing to set the 1st, 3rd and 4th entries in `c` to 0. Check it.
10. Use *boolean* fancy indexing to set the 1st and last entries in `a` to -1. Check it.

<go over solutions>